

## Identifying Key Entities in Recipe Data: Report

## Problem Statement

The lack of tagging in recipe data restricts the user experience on online cooking and meal-planning platforms. Recipes require manual annotation of ingredients, quantities, and titles which is time-consuming and inconsistent.

## Business Objective

The business objective is to leverage the increasing popularity of online cooking platforms and meal-planning apps by enhancing the user experience. This can be achieved by implementing a **custom-named entity recognition (NER)** model to automatically tag ingredients, quantities and recipe names. This automation will streamline the process of organising recipes, improve searchability and enable users to easily find recipes based on available ingredients, portion sizes or specific dietary requirements. This will ultimately reduce the labour-intensive and inefficient manual tagging process, providing a more accessible and efficient way for businesses in the food and recipe industry to manage their recipe databases.

## Dataset Description

This data set comprises culinary recipes with a focus on ingredient extraction and analysis. Each recipe features a structured ingredient list with labelled components, identifying ingredients, quantities and units. This diverse collection supports tasks such as understanding recipes and discovering culinary knowledge, enabling the development of models for information extraction in the culinary domain.

**Data Structure:**

The given data is in JSON format, representing a structured recipe ingredient list with Named Entity Recognition (NER) labels. Below is a breakdown of the data fields:

```
[  
    {  
  
        "input": "6 Karela Bitter Gourd Pavakkai Salt 1 Onion 3 tablespoon Gram flour besan 2 teaspoons Turmeric powder  
wder Haldi Red Chilli Cumin seeds Jeera Coriander Powder Dhania Amchur Dry Mango Sunflower Oil",  
        "pos": "quantity ingredient ingredient ingredient ingredient quantity ingredient quantity unit ingredi  
ent ingredient ingredient quantity unit ingredient ingredient ingredient ingredient ingredient ingredient ingre  
ent ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient"  
    },  
    {  
        "input": "2-  
1/2 cups rice cooked 3 tomatoes teaspoons BC Belle Bhat powder 1 teaspoon chickpea lentils 1/2 cumin seeds  
white urad dal mustard green chilli dry red 2 cashew or peanuts 1-1/2 tablespoon oil asafoetida",  
        "pos": "quantity unit ingredient ingredient quantity ingredient unit ingredient ingredient ingredient ingredie  
nt quantity unit ingredient ingredient quantity ingredient ingredient ingredient ingredient ingredient ingredien  
t ingredient ingredient ingredient ingredient quantity ingredient ingredient ingredient ingredient quantity unit ingredie  
nent"  
    }  
]
```

Key	Description
input	Contains a raw ingredient list from a recipe.
pos	Represents the corresponding part-of-speech (POS) tags or NER labels, identifying quantities, ingredients, and units.

## Data Ingestion and Preparation

## Identifying Key Entities in Recipe Data: Report

Shape: 285 rows and 2 columns

### Information:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 285 entries, 0 to 284
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  --
 0   input   285 non-null    object
 1   pos     285 non-null    object
dtypes: object(2)
memory usage: 4.6+ KB
```

### Recipe Data Manipulation and Validation:

- Split the input and pos columns to create the tokenized columns for both i.e. input\_token, pos\_token.

Sample below:

	input	pos	input_token	pos_token
0	6 Karela Bitter Gourd Pavakkai Salt 1 Onion 3 tablespoon Gram flour besan 2 teaspoons Turmeric powder Haldi Red Chilli Cumin seeds Jeera Coriander Powder Dhania Amchur Dry Mango Sunflower Oil	quantity ingredient ingredient ingredient ingredient ingredient quantity ingredient quantity unit ingredient ingredient ingredient quantity unit ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient ingredient	[6, Karela, Bitter, Gourd, Pavakkai, Salt, 1, Onion, 3, tablespoon, Gram, flour, besan, 2, teaspoons, Turmeric, powder, Haldi, Red, Chilli, Cumin, seeds, Jeera, Coriander, Powder, Dhania, Amchur, Dry, Mango, Sunflower, Oil]	[quantity, ingredient, ingredient, ingredient, ingredient, ingredient, quantity, ingredient, quantity, unit, ingredient, ingredient, ingredient, quantity, unit, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient, ingredient]

- Checked the length of the input\_token and pos\_tokens to make sure the lengths are equal. From the equality check, we have the following 5 rows with unequal length: [17,27,79,164,207]
- Checked the unique values in the pos column. The results were {'ingredient', 'quantity', 'unit'}
- Removed the rows with unequal length of input and pos tokens. Post removal, checked the shape of the dataset which was: (280, 6).
- Updated the input\_length and pos\_length columns to reflect the updated length i.e. 280, 6
- Checked the lengths again to find out if there is any more records of unequal length. Count of rows which are not of equal length :0

### Training and Validation data split

- The dataset is split into training and validation sets using a 70:30 ratio using train\_test\_split with a random\_state of 42.
- The output data frames are train\_df, val\_df. Checked records for both data frames
- Extracted X\_train, X\_val, y\_train and y\_val by extracting the list of input\_tokens and pos\_tokens from train\_df and val\_df and checked their length.
  - X\_train.shape : (196,)
  - X\_val.shape : (84,)
  - y\_train.shape : (196,)
  - y\_val.shape : (84,)
- Checked for unique labels in train\_df which was {'ingredient', 'quantity', 'unit'}

### Exploratory Data Analysis on Training Dataset

- Extracted the input tokens and its pos tags in training dataset and flattened it.
- Checked the length of input and pos tokens which were equal.
- Categorised tokens into ingredients, units and quantities by using extracted token function get a list of ingredients, units and quantities.

## Identifying Key Entities in Recipe Data: Report


- Created and used a function to get the top ingredients.

```
[('powder', 183),  
 ('Salt', 149),  
 ('seeds', 125),  
 ('Oil', 122),  
 ('Red', 120),  
 ('chopped', 115),  
 ('Green', 115),  
 ('Chilli', 113),  
 ('Coriander', 98),  
 ('Leaves', 94)]
```

- Created and used a function to get the top recipes used in recipes.


```
[('teaspoon', 221),  
 ('cup', 193),  
 ('tablespoon', 131),  
 ('tablespoons', 93),  
 ('grams', 82),  
 ('cups', 74),  
 ('inch', 72),  
 ('sprig', 62),  
 ('teaspoons', 57),  
 ('cloves', 55)]
```

- Plotted the top 10 frequently used ingredients from training dataset:



- Plotted the top 10 frequently used ingredients from training dataset:

## Identifying Key Entities in Recipe Data: Report



### Feature Extraction for CRF Model

- Defined keywords for unit and quantity and created a quantity pattern to work on fractions, numbers and decimals.
  - `quantity_pattern = re.compile(r'^\d+([/-]\d+)?(\.\d+)?$|^d+/\d+$|^d+\.\d+$|^d+$')`
- Loaded the spacy model -> `spacy.load("en_core_web_sm")`
- Defined a feature function for processing each token in the sentence.
- Converted the `X_train`, `X_val`, `y_train` and `y_val` into train and validation feature sets and labels
- Checked the length of training and validation features and labels, Output -> 196 and 84 respectively.
- Created **label\_counts** to count the frequencies of labels present in training dataset `y_train_flat` (flattened `y_train`) and retrieved the total samples.
  - Label count -> Counter ({'ingredient': 5323, 'quantity': 980, 'unit': 811})
  - Total samples -> 7114
- Compute class weights (inverse frequency method) by considering `total_samples` and `label_counts`

```
{'quantity': 2.419727891156463,  
'unit': 2.923962186600904,  
'ingredient': 0.44548813325818776}
```

- Applied penalizing factor on the ingredient label.
- Created **`X_train_weighted_features`** and **`X_val_weighted_features`** for extracting training and validation features along with their weights by using a method created for this purpose.

### Model Building and Training

- A CRF model was implemented using the `sklearn_crfsuite` library.
- CRFs are ideal for sequence labelling because they account for contextual dependencies between adjacent labels.

```
CRF(algorithm='lbfgs', all_possible_transitions=True, c1=0.5, c2=1.0,  
    max_iterations=100)
```

### Model Evaluation using CRF

- Evaluate on training dataset using
  - CRF by using flat classification report and
  - confusion matrix
- Classification reports were generated using training data set and training data

# Identifying Key Entities in Recipe Data: Report


Classification Report on Training Set:

	precision	recall	f1-score	support
ingredient	0.99	1.00	0.99	5323
quantity	1.00	0.98	0.99	980
unit	0.98	0.98	0.98	811
accuracy			0.99	7114
macro avg	0.99	0.99	0.99	7114
weighted avg	0.99	0.99	0.99	7114

Flat Classification Report on Training Data:

	precision	recall	f1-score	support
quantity	0.9969	0.9847	0.9908	980
unit	0.9802	0.9753	0.9778	811
ingredient	0.9934	0.9964	0.9949	5323
accuracy			0.9924	7114
macro avg	0.9902	0.9855	0.9878	7114
weighted avg	0.9924	0.9924	0.9924	7114

- Generated confusion matrix for training for the training data



- Model was saved as crf\_model.pkl

## Prediction and Model Evaluation


- Similar to training data, used CRF model to generate classification report on validation data.

Flat Classification Report on Validation Data:

	precision	recall	f1-score	support
quantity	0.9902	0.9878	0.9890	411
unit	0.9765	0.9302	0.9528	358
ingredient	0.9878	0.9962	0.9920	2107
accuracy			0.9868	2876
macro avg	0.9848	0.9714	0.9779	2876
weighted avg	0.9867	0.9868	0.9867	2876

- Generated confusion matrix using validation data

Identifying Key Entities in Recipe Data: Report



Error analysis on validation data

- Performed error analysis on validation data to find out mis-classified samples.

Total validation samples: 2876  
Total errors found: 38  
First 5 errors: [('ingredient', 'unit'), ('unit', 'ingredient'), ('unit', 'ingredient'), ('quantity', 'ingredient'), ('quantity', 'ingredient')]

- Iterated through validation data (X\_val, y\_val\_labels, y\_pred\_val) and compared true vs. predicted labels. Collected error details, including surrounding context, previous/next tokens, and class weights.

Total errors collected: 38  
Sample error details:  
{'token': 'cloves', 'true\_label': 'ingredient', 'predicted\_label': 'unit', 'prev\_token': '3', 'next\_token': 'garlic', 'class\_weight': 0.04454881332581878}  
{'token': 'Spoon', 'true\_label': 'unit', 'predicted\_label': 'ingredient', 'prev\_token': 'big', 'next\_token': 'oil', 'class\_weight': 2.923962186600904}  
{'token': 'cloves', 'true\_label': 'unit', 'predicted\_label': 'ingredient', 'prev\_token': 'seeds', 'next\_token': 'garlic', 'class\_weight': 2.923962186600904}  
{'token': 'is', 'true\_label': 'quantity', 'predicted\_label': 'ingredient', 'prev\_token': 'Pur', 'next\_token': '2', 'class\_weight': 2.419727891156463}  
{'token': 'to', 'true\_label': 'quantity', 'predicted\_label': 'ingredient', 'prev\_token': 'sugar', 'next\_token': 'tablespoons', 'class\_weight': 2.419727891156463}

- Changed error\_data into dataframe and then used it to illustrate the overall accuracy of validation data

Validation Error DataFrame:

	token	true_label	predicted_label	prev_token	next_token	class_weight
0	cloves	ingredient	unit	3	garlic	0.044549
1	Spoon	unit	ingredient	big	oil	2.923962
2	cloves	unit	ingredient	seeds	garlic	2.923962
3	is	quantity	ingredient	Pur	2	2.419728
4	to	quantity	ingredient	sugar	tablespoons	2.419728

Overall Accuracy on Validation Data: 0.9868

- Analysed errors found in the validation data by each label and displayed their class weights along with accuracy and also display the error dataframe with token, previous token, next token, true label, predicted label and context

## Identifying Key Entities in Recipe Data: Report



Error Analysis by Label:

Label	Errors	Total	Accuracy	Class Weight
quantity	5	411	0.9878	2.4197
unit	25	358	0.9302	2.9240
ingredient	8	2107	0.9962	0.0445

### Insights from validation dataset

- From the classification report we can draw the below insight:
  - Model performance is excellent with an overall accuracy of **98.68%** and **macro F1-score of 97.79%** indicates that your model is highly reliable across all three entity types
  - Ingredient entity type performs the best with **Precision: 0.9878, Recall: 0.9962, F1: 0.9920**. The model is quite confident and accurate in recognizing the **ingredients**.
  - **Quantity** entity also performs very well with F1-score of **0.9890**.
  - **Unit** has low score, still very strong performance, but relatively lower **recall** (0.9302). There are chances of mis-classifying the units compared to other entities.
- From the confusion matrix we have the following insights
  - Model accuracy is very high. The **diagonal values (406, 333, 2099)** represent correct predictions.
  - Almost all tokens are classified correctly across all three classes.
  - Similar to classification report, here also **ingredients** perform best where out of 2107 (2+6+2099) only 9 are mis-classified.
  - **Quantity** class is very clean where out of 411, only **5 misclassifications** (2 as unit, 3 as ingredient).
  - **Unit** has major mis-classification where **23 unit tokens** misclassified as ingredients (most significant confusion).

### Recommendations

No major recommendations, model looks very good. Units can be looked into a bit for further accuracy.

### Conclusion

- The given dataset was overall ok but few data quality issues were identified with 5 records where the input token and corresponding pos tag lengths were not matching.
- The primary entities extracted were 'quantity', 'ingredient', and 'unit'.
- Overall, a very good model with more than 98.68% accuracy and weighted F1-score 98.67%.
- Model is not overfitted as it works very well on unseen data.
- Due to its high-accuracy, The CRF model can reliably replace **manual tagging** of recipe data, saving time and reducing error.